
Native T-SQL Language Service

Analysis, Metadata, Features, and SQL Scripting

Design, implementation, and validation review

For [vscode-mssql PR #22860](#)

Central question. How can `vscode-mssql` provide deterministic, low-latency T-SQL completion, diagnostics, hover, signature help, definition, folding, document symbols, and script generation over one immutable metadata generation—while remaining useful with partial or offline metadata, avoiding network I/O on an interactive request, and staying invisible until a later private-preview consumer explicitly composes it?

Assessment at the reviewed head. PR #22860 establishes a coherent source-only language and scripting library. Its contracts are VS Code-neutral; requests pin one immutable metadata view; completion and hover can request background hydration without awaiting it; and diagnostics distinguish syntax truth from metadata-dependent claims. The branch adds no commands, settings, provider registration, controller, activation path, or extension bundle input. Review commit `b66db2d8f` closes the reproduced ordinary-SQL trust defects and the material Tier-B/C findings with focused regression coverage. Exact-head local validation passed 707 owned language/scripting tests, 47 adjacent service/observability tests, the 4,742-test broad MSSQL lane, SQL Projects, build, lint, formatting, hooks, and online packaging. All 26 initial threads plus one follow-up have explicit dispositions; a six-minute delayed audit found zero unresolved threads. The remaining gates are hosted CI completion and human review, not known semantic blockers.

Prepared for:	Karl Burtram
Primary change:	<code>dev/karl1b/query_05_language_service</code> at <code>b66db2d8f</code>
Extraction base:	<code>95cf790e5</code> ; 59 files, +26,391 / -0
Live PR base observed:	<code>53b707bbe</code> ; mergeable, human review required
Metadata prerequisite:	merged PR #22836 / shared metadata service in <code>main</code>
Report source parent:	<code>617567f17</code> in <code>kburtram/notes main</code>
Report snapshot:	31 August 2026, America/Los_Angeles

This is an as-built report. Historical refactor documents explain intent, but the pinned implementation, tests, PR state, and package artifacts control when the documents and code differ.

Contents

How to read this report	1
1 Executive summary	1
2 Placement in the refactor train	2
3 Public contracts and dependency boundaries	3
4 Document analysis pipeline	4
5 Metadata contract and honest degradation	5
6 Language feature behavior	6
7 Scheduling, routing, and failure containment	7
8 Definition and SQL scripting	8
9 Observability and privacy	9
10 Test design and validation	11
11 Known limitations and review closure	13
12 Integration and graduation plan	16
13 Final assessment	17
Evidence and source map	19

How to read this report

The report keeps implementation, intended invariants, executable evidence, and future integration separate:

Implemented

Present in the pinned PR head and traced to source.

Target invariant

A behavior or safety contract the library is intended to preserve.

Validated Exercised by a named test, build lane, package check, or observed repository state.

Open finding

An evidence gap, preview-graduation requirement, or limitation that remains at report freeze.

Future contract

A consumer, bridge, activation path, or maturity step not shipped by this PR.

Key takeaway

PR #22860 is not a replacement language service that users can turn on. It is the independently reviewable library beneath such a replacement: public data contracts, a tolerant analysis pipeline, native features, a pinned metadata adapter, strict scripting, routing/scheduling policies, observability, and tests. A later PR must supply a gated consumer and VS Code adaptation before any product behavior changes.

Evidence boundary

The strongest evidence is deterministic unit behavior and source-boundary inspection. Local packaging proves the tree packages; GitHub's package comparison more directly proves that the new source contributes zero size to the shipped VSIX at this head. Neither proves live-catalog semantic equivalence across the full T-SQL grammar. That is a later preview-graduation gate.

1 Executive summary

1.1 What the PR establishes

The change adds 37 production files under `src/sqlLanguage`, six under `src/sqlScripting`, fifteen focused unit files, and one observability classification edit. The architecture has four load-bearing properties:

1. **Pure request contracts and analysis.** Positions are zero-based UTF-16; payloads are JSON-safe; the core has no VS Code document or provider types.
2. **One truth generation per request.** Language work pins one immutable metadata view and reads readiness, environment, names, columns, parameters, keys, and definitions through that generation-stable seam.
3. **Honest partial behavior.** Absence and unavailability are different. The engine suppresses metadata-dependent claims, returns incomplete completion where appropriate, and preserves lexical/structural diagnostics when catalog validation is unavailable.

4. **Source-only private-preview boundary.** There is no package contribution, command, setting, activation event, controller, provider registration, Query Studio bridge, or AI consumer. Existing users cannot reach the code from this PR.

1.2 What the PR does not establish

It does not claim full T-SQL parsing, ScriptDOM equivalence, semantic-token or document-highlight support, live runtime integration, a production latency service-level objective, or a user-visible migration from the current SQL Tools Service language providers. The router includes a lazy bridge seam, but the bridge itself and the host composition that chooses it are deferred.

Private-preview meaning in this PR

No dedicated setting is added because there is no activated feature to gate. That is safer than a dead setting that suggests usable behavior. The first integration PR must require the umbrella experimental/private-preview flag *and* a dedicated feature-area flag, default both off, hide all contributed UI, and register runtime consumers only when the effective gate is true at activation.

2 Placement in the refactor train

PR #22860 is extraction PR05. Perfctest, core observability, SQL Data Plane/STS2 foundation, and the metadata substrate are already in **main**; the Debug Console remains an independent UX track. PR05 builds directly from **main**, not from Debug Console PR03, and consumes only the accepted metadata APIs.




Figure 1: PR05 placement. The language library and metadata adapter exist, but every path from a user-visible consumer is still a future, gated composition step.

This placement is why the GitHub package comparison reports exactly the same MSSQL VSIX size for target and PR branches: 85,347 KiB in both. The extension bundler has no `sqlLanguage` or `sqlScripting` input at this head.

3 Public contracts and dependency boundaries

3.1 Consumer-neutral API

`sqlLanguage/api.ts` defines request and response data for completion, hover, signature help, diagnostics, definition, folding, document symbols, highlights, and semantic tokens. Every request carries text, version, URI, and position/range data without importing VS Code. This boundary lets a later host translate editor objects once, then test the engine with ordinary values.

The same file defines feature maturity and a capability table. The shipped native capability table marks completion, hover, signature, diagnostics, definition, folding, and document symbols as preview; highlights and semantic tokens remain off. Having types for a feature is therefore not the same as advertising an implementation.

Table 1: Native feature surface at the pinned head.

Feature	Capability	As-built behavior
Completion	Preview	Context/expectation-gated candidates, deterministic rank, partial-readiness honesty, lazy hydration kick.
Hover	Preview	Local, overlay, catalog, system, routine, built-in, and description facts when their readiness tier permits.
Signature help	Preview	Built-ins, catalog functions/procedures, nested calls, and EXEC argument tracking.
Diagnostics	Preview	T1 lexical/structural errors plus freshness-gated T2 208/207/209 binder warnings.
Definition	Preview	In-document local targets or virtual scripted catalog definitions with anchors and fidelity.
Folding	Preview	Batch, statement, comment, and region ranges.
Document symbols	Preview	Batch/statement structure and simple create/alter object declarations.
Highlights	Off	Interface reserved; native engine returns no result.
Semantic tokens	Off	Interface reserved; native engine returns no result.

3.2 Purity and import direction

Core analysis, feature computation, static data, provider interfaces, and scripting emitters are library code. Concrete host concerns live under `host`: metadata freshness, diagnostic journaling, timeout/yield behavior, and routing. The one sanctioned bridge from language metadata to the product catalog model is `provider/catalogProvider.ts`.

Design invariant

Interactive feature code consumes a synchronous `IPinnedMetadataView`; it does not hold a mutable store, lease, network client, VS Code document, or STS service. Definition text is the one deliberate async lookup on the pinned view, and strict scripting asks the host to establish freshness before it creates the engine.

4 Document analysis pipeline

4.1 From text to reusable structure

The native engine analyzes one immutable text snapshot, then reuses structural work across feature calls when document version and text length match. The stages are deliberately tolerant:

1. **Text snapshot** maps offsets to zero-based UTF-16 line/character positions and back, including ordinary LF/CRLF and mixed input.
2. **Lexer** recognizes T-SQL words, identifiers, punctuation, comments, strings, variables, SQLCMD-sensitive regions, Unicode horizontal whitespace, and `GO / GO n` line semantics. Review hardening closes the non-breaking-space and construct-opener-at-EOF coverage gaps.
3. **Segmenter** partitions batches and statements while respecting module bodies and rule-sensitive constructs.
4. **Sketch parser** extracts the structure features need: query scopes, sources, CTEs, select items, declarations, `EXEC`, DML targets, and table declarations. It remains total on incomplete editor text.
5. **Overlay** models script-local state such as temporary tables, declared table variables, scripted tables, and batch-scoped variables.
6. **Binder and cursor expectation** combine local structure, the effective database, and one pinned metadata view for a particular feature request.




Figure 2: As-built analysis pipeline. Syntax structure is cached per document analysis; metadata is pinned and binding is performed for each feature request so one response cannot mix generations.

4.2 Why a sketch instead of a full AST

Editor input is commonly unfinished. A full grammar can make a single missing token collapse useful context; an unbounded hand-written parser can accidentally become a second compiler. The sketch deliberately extracts only what current features consume and records uncertainty. The cursor-expectation layer improves completion category selection and suppresses suggestions in declarations, data-type positions, comments, strings, and other quiet contexts, but it is not the richer parser-v2/AST envisioned by later design notes.

Semantic ceiling

The tolerant pipeline is broad, not complete. Dynamic SQL, unusual procedural forms, deeply dialect-specific syntax, and ambiguous recovery can exceed the sketch. The correct failure mode is suppression, incomplete results, or fallback—never a confident catalog error manufactured from an untrusted parse.

5 Metadata contract and honest degradation

5.1 Generation-stable provider seam

`provider/types.ts` separates a mutable provider from the immutable view used by a request. The provider exposes current generation, environment, readiness, a pin operation, databases, an explicit hydration request, and a change event. The pin exposes synchronous resolution and projection over an opaque `LangObjectRef`; consumers cannot retain mutable catalog objects or invent cross-generation identity.

Readiness is sectioned across objects, columns, parameters, foreign keys, and definitions. The global mode—full, lite, partial, or offline—does not erase section truth. Resolution is a tagged result: resolved, default-schema choice, ambiguous, not found, or unavailable. This prevents the dangerous equation “not returned = does not exist.”

5.2 Adapter over the shared metadata service

`catalogProvider.ts` is the only production adapter from the merged metadata substrate to language types. It maps a compact immutable `CatalogSnapshot` into the pinned language surface, carries database/case/capability context, and returns honest unknown states when no snapshot exists.

When a feature discovers that object names are ready but columns are not, the adapter may request hydration. The current implementation coalesces to one in-flight refresh, avoids starting while already loading/stale, makes at most one kick per generation, and enforces a 30-second floor. The present metadata service refreshes the full ladder; the language seam is already shaped so later section-level hydration does not change feature contracts.




Figure 3: Lazy hydration without hot-path I/O. The current request finishes from generation g ; a later request observes $g + 1$.

5.3 System and offline views

The null provider supplies explicit offline readiness and no false catalog facts. Fixture providers simulate full, partial, and lazy-hydration states for deterministic tests. The system-catalog decorator supplies a curated names-only fallback for `sys` and `INFORMATION_SCHEMA`; live catalog facts win, identifiers are synthetic and generation-local, and the fallback does not claim column types, nullability, keys, definitions, or other facts it does not carry.

No network on the feature path

Completion, hover, signature help, and ordinary diagnostics read the pin synchronously. A hydration request is advisory and asynchronous. Strict scripting is different by design: its host establishes an explicit freshness verdict before constructing the pinned engine. The separation makes latency and truth policy visible rather than allowing an accidental catalog call inside fuzzy matching or binding.

6 Language feature behavior

6.1 Completion

Completion first classifies cursor expectation, then uses the legacy context classifier and bound statement facts to choose producers. Producers cover statement keywords/defaults/snippets, member access, table sources, joins and foreign-key predicates, expressions and columns, variables, built-ins, star expansion, `INSERT`, `UPDATE`, `EXEC`, `DECLARE`, `USE`, data types, and databases.

Candidate selection is deterministic: category gates run before fuzzy matching; scores are stable; ordinal comparison breaks ties; output is capped at 600. Partial metadata does not produce an authoritative empty list. Member-access on an object with unloaded columns returns an incomplete reason and can kick background hydration.

At the pinned head, the architecture and contexts follow that policy, including the review-repaired statement-location fallback, derived aliases, `UPDATE TOP`, `db..table`, temp overlays, case-sensitive dedupe, and bare-`alias`. edits. Their regression cases and dispositions are summarized in [section 11.2](#) and [table 6](#); “preview” is not used to excuse false-positive behavior.

6.2 Diagnostics

Diagnostics intentionally separates two truth tiers:

- **T1 lexical/structural diagnostics** come from the text, tokens, segments, and tolerant statement structure. They include lexical faults, malformed `GO`, bracket/parenthesis/statement structure, unrecognized statements, and recoverable `SELECT` syntax.
- **T2 binder diagnostics** depend on catalog trust: invalid object 208, invalid column 207, and ambiguous column 209. A host freshness verdict of `notValidated`, partial readiness, dynamic/unsupported syntax, uncertain database context, or other recorded suppression reason prevents these warnings.

The diagnostic result records suppression counts/reasons so “quiet because valid” can be distinguished from “quiet because the engine lacked authority.” Whole-document work is resumable by statement/work unit for scheduler time slicing.

The 62-case honesty suite validates the constructs it contains, not all legal T-SQL. [Section 11.2](#) records newly reproduced false positives for alias assignment, cursor syntax, `WITHIN GROUP`, `AT TIME ZONE`, multi-column derived aliases, pasted NBSP, `UPDATE TOP`, and default-schema multipart names. These are corpus omissions and implementation defects, not evidence against the T1/T2 separation itself.

6.3 Hover and signature help

Hover resolves schema/database/name-chain parts, aliases, projected aliases, CTEs, derived tables, temporary/script tables, table variables, variables and parameters, catalog routines, and built-ins. H7 descriptions

appear only when description readiness is true. Cross-database and **USE** handling share the database-context map used by binding.

Signature help scans the enclosing call without requiring a complete parse. It understands nested and grouped built-in calls, catalog functions/procedures, and **EXEC** argument positions. Both features stay quiet in comments, strings, SQLCMD regions, and positions where the scanner cannot make an honest call.

6.4 Definition, folding, and symbols

Definition has two destinations. Local aliases, variables, CTEs, temporary objects, table variables, and derived columns return in-document ranges. Catalog objects or columns route through the scripting service and return virtual content plus semantic anchors. Folding and document symbols use only the analysis snapshot and do not require catalog I/O.

Table 2: Feature truth dependencies and degradation.

Feature	Primary truth input	When truth is incomplete
Completion	Cursor expectation, sketch, overlay, objects/columns/keys	Limit categories, return incomplete, request hydration; retain static/built-in/local candidates.
Diagnostics	Tokens/structure; then bound objects/columns	Preserve T1; suppress affected T2 and account for the reason.
Hover	Local structure or specific catalog sections	Return only facts whose section is ready; no invented type/detail.
Signature	Call scanner, built-ins, routine parameters	Built-ins/local syntax remain; catalog signatures require parameter readiness.
Definition	Local declarations or scriptable catalog target	Local range when known; virtual script with provenance; unavailable/refusal when definition truth is insufficient.
Folding/symbols	Snapshot, tokens, segments, sketch	Continue structurally; no metadata dependency.

7 Scheduling, routing, and failure containment

7.1 Sliced diagnostics scheduler

`host/scheduler.ts` is VS Code-free. A document or metadata change restarts a default 300 ms debounce. The pass then performs work in default 8 ms slices and yields by macrotask between slices. Its stamp normally combines document version and metadata generation; the scheduler rechecks the stamp before starting, after an asynchronous pass factory returns, between slices, and before publishing.




Figure 4: Diagnostics scheduling. Old results are abandoned instead of being published after either document edits or metadata-generation drift.

7.2 Feature router

`host/router.ts` centralizes rollout and fallback. A consumer supplies an engine preference, capability table, rollout threshold, lazy bridge factory, and optional diagnostics observer. Native work is eligible only when preference and maturity allow it. After two consecutive thrown/rejected failures for a feature, its circuit opens and later calls go directly to fallback until reset. Failures are isolated per feature, so diagnostics cannot disable completion.

The router also wraps native work in a default five-second `Promise.race`. That is effective only when the native operation yields. Completion, hover, signature, pull-path diagnostics, folding, and symbols currently compute synchronously, so a timer cannot preempt a pathological call that blocks the extension-host event loop; only definition and other genuinely asynchronous paths can observe the timeout while in flight. Three core tests cover preference/maturity routing and throw-driven circuit/reset behavior, but no test proves timeout preemption. The scheduler’s sliced diagnostics path remains the real cooperative bound for hosted whole-document diagnostics.




Figure 5: Per-feature routing. The native library can be introduced incrementally while retaining a lazy fallback and containing repeated failures. The bridge and composition are not part of PR05.

8 Definition and SQL scripting

8.1 Scripting contract

`sqlScripting/api.ts` defines create, alter, create-or-alter, drop, drop-and-create, select-top, insert, update, delete, and execute operations. Results carry script text, anchors, fidelity, notes, unavailability reason, and metadata provenance.

Fidelity is explicit:

- **F2** is definition-level output: a stored module definition or a table statement with the required detailed metadata.
- **F1** is structurally useful but incomplete output, such as a table definition degraded by missing optional details.
- **F0** is a template, notably DML scaffolding, and is never presented as a faithful reconstruction.

Module scripting uses token-level head rewriting so create/alter/create-or-alter changes do not replace keywords inside comments, strings, or bodies. Table scripting synthesizes columns, types, defaults/computed expressions, identity, and constraints only when corresponding metadata is ready. DML emitters produce honest templates. The writer records semantic anchors so definition navigation can land on the object, parameter, column, or statement rather than merely opening line zero.

8.2 Strict freshness flow

The concrete scripting host calls the metadata lease’s strict policy before it pins. A host-side timeout is a backstop around the metadata policy timeout. Rejection, timeout, or an unavailable strict verdict refuses the operation; explicit offline-snapshot policy can instead allow cached output with provenance and an honest banner. This makes scripting stricter than interactive completion because generated DDL is more likely to be copied and executed.




Figure 6: Definition and scripting truth flow. Local navigation stays in the document; catalog navigation produces a provenance-bearing virtual script or an explicit refusal.

9 Observability and privacy

The engine emits bounded spans for lexing, segmentation, parsing/sketching, completion, hover, signature help, diagnostics, definition, routing, and scripting. `perfTelemetry.ts` classifies `sqlLanguage.*` and `sqlScripting.*` under the existing `sqlLanguage` feature family; the accepted observability registry already covers those prefixes.

Normal diagnostics use safe scalar measures such as duration, counts, route, readiness, cache/hydration state, suppression reasons, and result shape. Rich completion diagnostics can classify prefix/name-part/top-label values as `user.text` or `object.name`; the shared diagnostics layer digests these by default and exposes them only under the explicitly elevated, time-bounded capture policy. The privacy claim is therefore not “identifiers never enter an event object.” It is that sensitive fields are classified at the call site and governed by the accepted capture/redaction boundary.

Observability contract

Do not add raw document text, SQL batches, connection strings, result rows, or unclassified object names to a language span. New attributes must be registered and typed before emission. Performance evidence should use counts, durations, readiness and suppression state; diagnosis that truly needs content must remain an explicit elevated-capture operation.

10 Test design and validation

10.1 Focused semantic corpus

The 15 owned language and scripting unit files execute 707 tests at the pinned head: 707 passed, zero failed, zero skipped. The value is not just volume; the suites target incomplete syntax, partial truth, deterministic ordering, stale publication, refusal behavior, and every material review repair. The exact-head sign-off also ran 47 adjacent scripting-service and observability tests, for 754 focused passes in total.

Table 3: Focused language and scripting inventory.

Area	Tests	Principal invariants
Completion	82	Member access, table sources, joins/FKs, expressions, CTE/derived/temp objects, DML/EXEC/DECLARE/USE, DDL declarations, star expansion, case-aware dedupe/edit ranges, partial-readiness honesty, latency budget, statement starts, static system catalog.
Cursor expectation	17	Category/confidence selection and quiet positions before completion producers run.
Core / router	24	Lexer coverage/Unicode whitespace, segmenter, feature routing, async timeout, maturity, fallback/circuit behavior, and LS-0 surface.
Definition	56	Aliases, variables, CTE/temp/script-local objects, catalog table/module scripts, honesty ladder, and capability routing.
Diagnostics	181	Lexical, G0, structural and recovery T1; 208/207/209 T2; expanded zero-unexpected-diagnostic honesty corpus; overlay/case/cache suppression; lazy hydration; engine pass.
Diagnostics freshness	6	Freshness gate and scheduler drift classification.
Hover	124	Catalog and local symbol families, H7 descriptions, partial truth, USE/cross-database/MERGE, DML positions, never-hover contexts, and system catalog.
Journal	3	Normal versus elevated-capture diagnostic fields.
Providers	16	Fourslash fixture parsing, fixture behavior, real CatalogBuilder/snapshot mapping, large filtered-prefix search, and unknown-case/offline honesty.
Scheduler	9	Debounce, sliced execution, stale document cancellation, metadata drift, and asynchronous factory recheck.
Signature	60	Built-ins, nesting/grouping, catalog routines, EXEC, and honesty positions.
Sketch	25	Total structural modeling across supported statement families, multipart names, module overlays, modern expression grammar, and incomplete text.
System catalog	16	Exact TVF result shapes, static data helpers, live-over-fallback precedence, and diagnostic interaction.
SQL scripting	77	Module rewrites/anchors/honesty; table F2/F1/anchors; comment-injection safety; DML; capabilities/routing; provenance/offline banners.
Strict scripting host	11	Freshness policy, timeout/refusal, pin ordering, provenance, and explicit offline behavior.

Area	Tests	Principal invariants
Total	707	Zero failures and zero skips in the owned focused selection.

10.2 Broader local validation

Validation was rerun after review commit `b66db2d8f` on the topic branch rooted at extraction base `95cf790e5`:

Table 4: Independent local validation at `b66db2d8f`.

Lane	Result	Interpretation / exclusion
MSSQL build	Pass	Toolkit plus MSSQL extension typecheck, emit, extension/webview/notebook bundles.
MSSQL lint / repository format	Pass	Target ESLint and full-tree Prettier check completed without errors.
Owned language/scripting	707 pass; 0 fail	All files in table 3 ; no skipped cases.
Adjacent service/observability	47 pass; 0 fail	Scripting service and three accepted observability suites; 754 total focused passes.
Broad MSSQL hosted tests	4,742 pass; 0 fail; 10 skip	Inverse selection excludes only two independently documented current-main Windows baseline suites: <code>Utility Tests - Path handling</code> and <code>QueryCompletionSoundService</code> .
SQL Projects hosted tests	205 pass; 0 fail; 6 skip	Executed in its separate VS Code host.
Pre-commit / diff audit	Pass	Staged hooks and <code>git diff -check</code> clean; validation-generated toolkit locales restored; no manifest, setting, command, or lockfile change.
Production online package	Pass	Released STS <code>6.0.20260827.1</code> ; local artifact 1,542 files / 104.9 MB. Metafile contains no new language/scripting bundle input.

10.3 Review-machine scale and differential probes

The Fable review added useful **Reported evidence** evidence beyond the fixture suites. At pre-fix head `eafbbdc046`, it ran the then-current language/scripting and observability suites, hydrated the real `PR04 MetadataStore` with a 10,201-object B16 fixture, adapted it through `CatalogLanguageMetadataProvider`, and exercised `NativeSqlLanguageEngine`. It also ran malformed-input/depth probes and compared the curated system catalog with a live SQL Server 2025 instance. These observations motivated the repair corpus; exact-head local validation then reran the product suites at `b66db2d8f`.

Table 5: Pre-fix Fable review-machine measurements at `eafbbdc046`; observations, not product SLAs.

Probe	Observed	Interpretation
FROM completion, 10,201 objects	p50 4.0 ms; p95 8.7 ms	Current full-scan/sort cost is usable on the probe, but remains linear and uncached for empty-prefix catalog requests.
Member access / alias columns / hover	p50 about 0.1 ms; p95 ≤ 1 ms	Pinned lookup and local binding paths are fast in the fixture.
Completion at end of 2,000 statements	p50 0.8 ms; p95 35.6 ms	Tail latency shows why production size-tier budgets are still needed.
Diagnostics, 2,000 statements	68.5 ms cold; near 0 cached	Analysis reuse is effective when its cache key is sound.
Document symbols / folding, 2,000 lines	0.3 / 1.0 ms	Pure structural features are inexpensive on the probe.
Malformed/depth sweep	No throws; 5,000 parens in 13 ms	Supports total/tolerant direction; the separately found opener-at-EOF coverage hole is repaired and regression-tested.

The same live-catalog comparison found that all 157 curated built-in function names were genuine and identified two incorrect TVF result shapes. Review commit `b66db2d8f` removes the input handle names from `dm_exec_sql_text` / `dm_exec_query_plan`, pins their exact outputs, and adds a real `CatalogBuilder`/snapshot/provider-to-engine suite rather than relying only on fixture providers or hand-built handles.

10.4 GitHub validation snapshot

At report freeze, GitHub reports the PR open, mergeable, non-draft, labeled **Automated Review**, and still requiring human review. At final head `b66db2d8f`, build-and-test, actions and JavaScript/TypeScript analysis, CodeQL, normalization, CLA, package creation, and five of six Linux/macOS/Windows x64/arm64 VSIX-launch jobs are green. The Windows arm64 launch job and Azure VSCode MSSQL - Pull Request validation are still in progress/pending and are not represented as passed.

The GitHub package comparison is stronger than the absolute local size for this scope question: target and head are both 85,347 KiB for MSSQL, with zero delta for every packaged extension. Codecov reports 94.34331% patch coverage with 1,041 changed lines uncovered and project coverage rising from 78.22% to 79.86%. The largest uncovered concentration is the production metadata adapter; hover, diagnostics, definition, completion, sketch/context, cursor expectation, engine, and text snapshot also retain uncovered branches.

Evidence boundary

High patch coverage and 754 exact-head focused passes support the implementation, but the important properties are semantic: one generation per answer, suppression when metadata authority is unavailable, no stale diagnostic publication, deterministic ranking, and strict scripting refusal. Coverage percentage alone cannot prove these; the differential probes were valuable precisely because they found ordinary constructs missing from the prior honesty corpus. Conversely, zero bundle-size delta proves inert packaging, not live feature quality.

11 Known limitations and review closure

11.1 Review snapshot: 27 dispositioned threads

The pre-fix PR had 26 current review threads at `eafbbdc046`: two from Copilot and 24 from a later source-attributed Fable review posted through the PR author's account. Every initial thread received an explicit reply ending (`Source: GPT-5.6 Sol`) and was resolved after review commit `b66db2d8f`. One immediate Copilot follow-up questioned the documented newline precondition of `withHeader`; a complete production call-site audit showed that every caller supplies a terminating CRLF, so it was dispositioned and resolved without defensive churn. After a six-minute quiet period, the live API returned zero unresolved threads at the unchanged head.

The two original Copilot threads were narrow and are fixed:

1. `sqlScripting/scriptWriter.ts` now counts standalone CR, LF, and CRLF correctly in writer and prepended-header anchors, with mixed-ending regression coverage.
2. `sqlScripting/scriptingService.ts` now says synonym definition text is unavailable through this scripting engine rather than making an inaccurate hydration claim.

11.2 Tier A: ordinary-SQL trust defects

The eight Tier-A findings were not architectural failures, but they were the exact user-trust failure class the diagnostic honesty design is intended to prevent. All eight are repaired at `b66db2d8f` with focused zero-unexpected-diagnostic or completion-suppression coverage.

- A1. Trailing comments and unterminated constructs** (`nativeEngine.ts`). Lexical suppression now runs before the undefined-statement fallback; comment, string, SQLCMD, and exact-EOF opener cases return no completion.
- A2. Alias assignment and positioned update** (`diagnostics.ts`). Select-list `alias = expr` heads and `CURRENT OF` cursor names are excluded from column diagnostics.
- A3. WITHIN GROUP (ORDER BY ...)** (`diagnostics / keyword data`). Contextual `WITHIN` plus ordered-set clause handling prevents the false missing-BY recovery.
- A4. AT TIME ZONE** (`diagnostics / keyword data`). Contextual `AT/ZONE` entries break the false expression-unit run without reserving legal identifiers.
- A5. Derived/VALUES alias list $v(a,b)$** (`sketch/index.ts`). Alias column lists are balance-skipped so their commas cannot create phantom sources.
- A6. Non-breaking space in pasted SQL** (`lexer.ts`). Unicode horizontal space separators, including `U+00A0`, are whitespace rather than identifier characters.
- A7. UPDATE TOP (n) target** (`sketch/index.ts`). Update sketching skips TOP count forms and optional PERCENT before reading the target.
- A8. Default-schema shorthand `db..table`** (`sketch/index.ts`). Empty multipart components are preserved, and binder arity uses raw parts to reach honest cross-database/linked-server handling.

11.3 Tier B/C: correctness, safety, and evidence hardening

The other 16 Fable threads were triaged in five groups. Material correctness and safety issues were fixed; two design/optimization suggestions received narrower evidence-backed dispositions:

Table 6: Source-attributed review findings and final dispositions.

Group	Findings	Disposition
Parser / overlay honesty	IS NOT DISTINCT FROM; labels/synonyms; script-created modules and qualified objects; EOF coverage; GO corpus.	Repaired with grammar, overlay, lexical-coverage, and table-driven regression tests.
Completion / provider	Temp DML targets; case-aware dedupe; bare-alias edits; filtered prefix cap; unknown case rule; empty-prefix scan.	Correctness paths repaired. Prefix search grows past 1,200 filtered-out items. Empty-prefix caching deferred: measured p95 is 8.7 ms and cache invalidation complexity is not justified without a budget regression.
Host correctness	Same-version/same-length cache collision; synchronous router timeout semantics.	Analysis and diagnostic memos require exact text. Async timeout is tested and cooperative semantics documented; synchronous interactive work remains self-bounding and diagnostic publication remains sliced.
Scripting safety	Comment placeholders/refusals can interpolate pathological identifiers or type text.	Reused <code>sanitizeCommentText</code> across DML and refusal paths; adversarial comment-terminator tests pass.
Static data / integration evidence	Two <code>dm_exec_*</code> TVF result schemas; fixture-only production adapter coverage.	Corrected exact TVF shapes and added real builder/snapshot/provider-to-engine coverage.

These findings are resolved or explicitly dispositioned in the PR. The expanded corpus now pins the reproduced cases; it still does not claim complete T-SQL grammar equivalence or replace preview-stage real-server validation.

11.4 Deliberate implementation limits

- The sketch is not a full T-SQL AST and the feature set is not a ScriptDOM equivalence claim.
- Highlights and semantic tokens are off.
- The bridge/fallback adapter is a router seam, not an implementation in this PR.
- The metadata adapter’s lazy kick currently refreshes the full metadata ladder; section-specific hydration remains a future optimization.
- The system fallback is names-only and deliberately cannot support type/detail/key/definition claims.
- No live editor consumer, settings hierarchy, provider registration, Query Studio integration, AI completion, SQL document service, sqlcmd preprocessing, execution batch splitter, or query execution helper ships here.
- Unit fixtures cover a wide semantic corpus and the review added one SQL Server 2025/static-data plus 10,201-object store probe, but broad real-server collation, permissions, dynamic SQL, uncommon grammar, Azure/serverless behavior, and remote extension-host performance remain preview validation work.

11.5 Branch-drift boundary

The topic includes `main` through extraction base `95cf790e5`. At report freeze the live PR API identifies a later base `53b707bbe`, and the locally fetched `origin/main` has advanced again to `e0f18fd8d`. GitHub currently calculates the PR as mergeable, but a later main refresh should repeat the full build/lint/focused/broad/package gates before the final merge decision.

12 Integration and graduation plan

12.1 Next implementation step

The next PR should be a gated consumer, not another implicit activation edit scattered across the extension. Before that consumer is made reachable, the eight Tier-A cases in [section 11.2](#) must be fixed and pinned by the honesty/suppression corpus. The consumer PR should then:

1. add one composition root that adapts VS Code documents and result types to the consumer-neutral contracts;
2. require `mssql.enableExperimentalFeatures` and a dedicated default-off native-language/AI feature flag before registration;
3. hide contributed commands and UI with complete manifest `when` clauses, not only disabled command enablement;
4. choose native versus existing behavior through the router, provide the real lazy bridge, and preserve per-feature fallback;
5. wire the metadata provider from the accepted shared catalog service and dispose leases/change subscriptions correctly;
6. carry perf test and scenario tests for activation-off invisibility, gated activation, latency, fallback, and stale-result suppression; and
7. avoid coupling AI completion to a global language-provider cutover—AI can be an early gated consumer of the same analysis/provider seam.

12.2 Graduation gates

Before moving from private preview toward default-on, validation should add:

1. a live-SQL truth corpus across supported versions, Azure variants, collations, permissions, synonyms, temporary/script objects, and cross-database contexts;
2. side-by-side native/STS result comparison with categorized differences rather than raw count parity;
3. interactive p50/p95/p99 budgets for completion/hover/signature and sliced whole-document diagnostics across document/catalog size tiers;
4. remote and web/virtualized extension-host CPU/memory measurements plus metadata hydration/retrigger behavior;
5. UI-level tests for flag hierarchy, command/menu invisibility, reload semantics, circuit-breaker fallback, and disabling the preview; and
6. an explicit privacy audit of elevated rich fields, journal/export behavior, and any AI consumer payload boundary.

Table 7: Mechanism-to-evidence traceability.

Mechanism	Current evidence	Next proof
Pure analysis contracts	Import shape, fixture engines, hosted unit tests	Consumer adapter tests in the first activation PR.
Pinned-generation metadata	Provider types and fixture/null/system tests; reviewer real-store probe	CI end-to-end real-snapshot suite plus forced generation change.
No hot-path network wait	Synchronous pin contract, lazy kick tests, and reviewer real-store timings	Instrumented product/perftest assertion under hydration.
Completion determinism	Category/rank/ordinal tests; known suppression, dedupe, overlay, and edit-range findings	Close A1/A5–A8 and Tier-B completion cases; cross-platform replay.
Diagnostic honesty	T1/T2 split, 62-case honesty suite, freshness/drift tests; eight new ordinary-SQL gaps	Close Tier A and add a real-server/native-vs-STS differential corpus.
Stale-result suppression	Deterministic scheduler stamp tests	VS Code integration under edit/hydration storms.
Strict scripting provenance	Freshness/refusal/offline and emitter tests	Live permissions/encryption/definition corpus.
Failure containment	Three preference/maturity/throw-circuit/reset tests; sliced scheduler tests	Cooperative bound for synchronous work, timeout test, real bridge/fallback telemetry.
Private-preview invisibility	No manifest/controller/bundle inputs; zero package delta	Manifest/runtime gate tests when activation is introduced.

13 Final assessment

Key takeaway

PR #22860 is a strong independent extraction boundary. It turns the earlier language-service design into a concrete, testable library without prematurely changing the product. The most important choices are sound: tolerant analysis over incomplete editor text; generation-pinned metadata; explicit readiness and suppression; deterministic ordering; sliced, stale-safe diagnostics; provenance-bearing scripting; and per-feature fallback containment.

The later review materially strengthened the correctness disposition. Eight reproduced ordinary-SQL trust defects now have direct repairs and regression cases; the broader cache-identity, completion-edit, provider, scripting-sanitization, static-data, and real-adapter gaps are closed. The router timeout is explicitly cooperative rather than falsely presented as preemptive for synchronous work, and the only deferred optimization—empty-prefix provider caching—has measured p95 evidence below the current budget. The branch also remains safe for existing users because it is neither bundled nor activated.

Decision statement

Recommended disposition at this snapshot: automated review closeout is complete for the inert source-only foundation at `b66db2d8f`. The exact-head local matrix is green, all 27 threads have explicit dispositions, and the delayed audit reports zero unresolved threads. Merge remains subject to the still-running hosted jobs and human review. A later consumer PR must independently satisfy the full private-preview umbrella/feature gates, invisible command/menu contribution, runtime fail-closed behavior, and integration-level scenario validation before exposing this engine through a provider,

Query Studio surface, or AI completion path.

Evidence and source map

Reproducibility note

This report is pinned to an active PR. Review threads, Actions conclusions, main-branch refs, coverage, and mergeability are observations at the stated snapshot. Re-fetch the head and live PR state after any code or base-branch change before reusing its final disposition.

Evidence class	Primary sources used
PR state	PR #22860 , head b66db2d8f , extraction base 95cf790e5 , live API base 53b707bbe , 27 dispositioned review threads / zero unresolved after the delayed audit, Copilot review, source-attributed Fable review and probe report, exact-head sign-off comment, hosted-check rollout, package-size bot, and Codecov report.
Public API / host	<code>extensions/mssql/src/sqlLanguage/api.ts</code> ; <code>host/nativeEngine.ts</code> ; <code>host/router.ts</code> ; <code>host/scheduler.ts</code> ; <code>host/scriptingHost.ts</code> .
Analysis core	<code>core/text/textSnapshot.ts</code> ; <code>core/lexer.ts</code> ; <code>core/segmenter.ts</code> ; <code>core/sketch/*</code> ; <code>core/parser/cursorExpectation.ts</code> ; <code>core/binder.ts</code> ; <code>core/overlay.ts</code> ; <code>database/name/fuzzy/quote</code> helpers.
Language features	<code>features/completion.ts</code> ; <code>diagnostics.ts</code> ; <code>hover.ts</code> ; <code>signatureHelp.ts</code> ; <code>definition.ts</code> ; <code>folding.ts</code> ; <code>documentSymbols.ts</code> ; generated <code>keyword/built-in/default/system</code> data.
Metadata seam	<code>provider/types.ts</code> ; <code>catalogProvider.ts</code> ; <code>systemCatalogView.ts</code> ; <code>nullProvider.ts</code> ; <code>fixtureProvider.ts</code> ; merged metadata-service implementation from PR #22836 .
SQL scripting	<code>extensions/mssql/src/sqlScripting/api.ts</code> ; <code>scriptingService.ts</code> ; <code>scriptWriter.ts</code> ; <code>module/table/DML</code> emitters; strict host under <code>sqlLanguage/host</code> .
Observability	<code>extensions/mssql/src/perf/perfTelemetry.ts</code> ; accepted diagnostics registry and privacy contracts already in <code>main</code> ; journal tests.
Tests	Fifteen <code>extensions/mssql/test/unit/sqlLanguage*.test.ts</code> and <code>sqlScripting*.test.ts</code> files; hosted JUnit XML; independent <code>build/lint/broad-suite/package</code> record summarized in table 4 .
Design intent	<code>vscode-ext-refactor/language-service-docs/05-tsql-language-service-design.md</code> ; <code>current_completions_parser_design.md</code> ; <code>PROGRESS.md</code> ; proposed diagnostics/completion designs. Code controls where historical/future design differs from the extraction.
Private-preview policy	<code>C:/repos/work2/resume.md</code> and the accepted PR04 preview-gating implementation; no PR05 manifest or activation additions.